

Vehicle Routing Problem with Time Windows

Tath Kanchanarin

July 17, 2025

Problem Statement

Vehicle Routing Problem with Time Windows (VRPTW) is a well-known combinatorial optimization problem, extending the classic Capacitated Vehicle Routing Problem (CVRP). It involves determining the optimal set of routes for a fleet of vehicles to serve a number of geographically distributed customers, each with a specific time interval — known as a time window — during which service must occur with service time at each customer node, and each vehicle have limited capacity.

This problem is highly relevant in real-world logistics, where customers and businesses often operate with strict delivery time windows, adding a significant layer of complexity to the standard vehicle routing problem. There are two common variants of the time window constraint: hard and soft. In the hard time window variant, each customer must be served strictly within the specified interval; violating it renders a solution infeasible. The soft time window variant allows violations at the cost of incurring a penalty cost in the objective function.

The objective of VRPTW is to minimize the total distance traveled by all vehicles while ensuring that each customer is served within their respective time window and that vehicle capacity constraints are respected. Some other formulations also treat the number of vehicles, minimizing cost for vehicle usage, and waiting time, incurring cost when waiting on nodes as a secondary objective but in this project, we assume an unlimited fleet size that does not require setup costs is provided and waiting time is permitted implicitly with no penalty cost. Thus, the only optimization criterion is the total Euclidean (L2) travel distance across all routes with only constraints of time window and capacity of each vehicle.

Mathematically, VRPTW is defined on a directed graph $G = (V, A)$, where V is the set of nodes and A is the set of arcs (we will use Distance Matrix instead but they are equivalent). The depot is represented by two nodes: 0 (start) and $n + 1$ (return), while the rest are customers.




Figure 1: VRPTW instance with depot and customer time windows [2]

Each nodes $i \in V$ has an associated time window $[a_i, b_i]$, where:

- a_i is the earliest time at which service can start,
- b_i is the latest time by which service must be completed,
- s_i is the service duration at node i .
- t_{ij} is the travel time between nodes i and j .

For simplicity, we assume that depot at node 0 and $n+1$ have $[a_0, b_0] = [a_n + 1, b_n + 1]$, where they are the earliest and latest time service can start at depot, with $s_0 = s_{n+1} = q_0 = q_{n+1} = 0$; servicing time. and load constraint is 0.

feasibility condition

A commonly used feasibility condition on the depot's time window is:

$$b_0 \geq \max_{i \in V \setminus \{0\}} \{ \max \{ a_0 + t_{0i}, a_i \} + s_i + t_{i,n+1} \}$$

This bit of constraint ensures the following:

- $a_0 + t_{0i}$: Earliest possible arrival at customer i , assuming you depart from the depot as early as possible.
- a_i : Earliest time you're allowed to serve customer i .

- $\max\{a_0 + t_{0i}, a_i\}$: Actual service start time at i , respecting both the travel time and time window.
- s_i : service duration at customer i .
- $t_{i,n+1}$: travel time to return from customer i to the end depot.

This condition checks whether or not each of the customers can be served within the time window individually, if any of the trip from any customers is more than b_0 , then there will not be solution that depot cannot serve all customer within the time window. Since Solomon already have a valid solution, we will ignore this preprocessing step.

Methodology and Implementation

IP formulation

The model that we will be using follows VRPTW1 [3]

Sets and Indices

- V – Set of all nodes, including the depot (0) and end depot ($n+1$)
- N – Set of customer nodes, i.e., $V \setminus \{0, n+1\}$
- A – Set of directed arcs between nodes
- K – Set of vehicles
- $\delta^+(i)$ – Set of nodes reachable directly from node i
- $\delta^-(i)$ – Set of nodes that can reach node i

Parameters

- c_{ij} – Distance or cost of traveling from node i to j
- t_{ij} – Travel time from i to j
- $[a_i, b_i]$ – Time window for service at node i
- s_i – Service duration at node i
- q_i – Demand of customer i
- Q – Capacity of each vehicle

Decision Variables

- $x_{ijk} \in \{0, 1\}$ – 1 if vehicle k travels from node i to j , 0 otherwise
- $T_{ik} \in \mathbb{R}_+$ – Time at which vehicle k starts service at node i

Decision Variable

$$x_{ijk} = \begin{cases} 1 & \text{if vehicle } k \text{ travels from customer } i \text{ to } j \\ 0 & \text{otherwise} \end{cases}$$

T_{ik} = arrival time of customer i for vehicle k

q_k = current load of vehicle k

$$\min \sum_{k \in K} \sum_{(i,j) \in A} c_{ij} x_{ijk} \quad (5.1)$$

$$\sum_{k \in K} \sum_{j \in \delta^+(i)} x_{ijk} = 1 \quad \forall i \in N \quad (5.2)$$

$$\sum_{j \in \delta^+(0)} x_{0jk} = 1 \quad \forall k \in K \quad (5.3)$$

$$\sum_{i \in \delta^-(j)} x_{ijk} - \sum_{i \in \delta^+(j)} x_{jik} = 0 \quad \forall k \in K, \forall j \in N \quad (5.4)$$

$$\sum_{i \in \delta^-(n+1)} x_{i,n+1,k} = 1 \quad \forall k \in K \quad (5.5)$$

$$x_{ijk} \cdot (T_{ik} + s_i + t_{ij} - T_{jk}) \leq 0 \quad \forall k \in K, (i,j) \in A \quad (5.6)$$

$$a_i \leq T_{ik} \leq b_i \quad \forall k \in K, i \in V \quad (5.7)$$

$$\sum_{i \in N} q_i \sum_{j \in \delta^+(i)} x_{ijk} \leq Q \quad \forall k \in K \quad (5.8)$$

$$x_{ijk} \in \{0, 1\} \quad \forall k \in K, (i,j) \in A \quad (5.9)$$

5.1 Objective function: Minimizes the total travel distance. Here, c_{ij} is the distance between customers i and j , and x_{ijk} is a binary decision variable indicating whether vehicle k travels from i to j .

5.2 Each customer is visited exactly once by one vehicle.

5.3 Each vehicle leaves the depot exactly once. (This may be relaxed later.)

5.4 Flow conservation: Ensures that the number of vehicles entering each customer node equals the number leaving it.

- 5.5** Each vehicle must return to the depot (duplicated as node $n + 1$).
- 5.6** Time consistency: Ensures that the arrival time at customer j is after the departure from i , accounting for service and travel time.
- 5.7** Time windows: Enforces that customer i is visited within its time window $[a_i, b_i]$.
- 5.8** Capacity constraint: Ensures that the total demand served by each vehicle does not exceed capacity Q .
- 5.9** Binary constraint: Each decision variable x_{ijk} is binary.

constraint 5.3, 5.4, 5.5 ensures subtour elimination and valid solution by ensuring that each vehicle start and at the depot, which is different from another formulation that uses Miller-Tucker-Zemlin (MTZ) constraints to eliminate subtours. This formulation is computationally worse than MTZ, but it is easier to implement and understand.

constraint 5.6 can be split into two subcases

- $x_{ijk} = 1$, then $T_{jk} \geq T_{ik} + s_i + t_{ij}$, making sure arrival at the next customer is after the service time and travel time from previous customer
- $x_{ijk} = 0$, then $0 \leq 0$ which is trivial

constraint 5.7 needs to be applied with indicator variable, since if $x_{ijk} = 0$, then T_{ik} can be any irrelevant dummy variable, but if $x_{ijk} = 1$, then T_{ik} must be within the time window.

our implementation follows the above formulation closely in python, supplemented with .ipynb file alongside the report.

Experiment

The experiment is conducted on a set of standard benchmark instances from the Solomon's dataset [1], which is consisted of 25/50/100 customer nodes, with specific flag of clustered(C), randomly(R), randomly clustered(RC) indicating the type of customer distribution on the plane. The experiment is conducted on macbook pro with m3 max chip and 32GB of RAM, under python 3.10.18 environment with CPLEX version: 22.1.1.0, licensed for academic use through python API docplex wrapped jupyter notebook with no thread limits and timeout limitations on solver.

The dataset is expected to be output in the format of " $|X|, |T|, |X| + |T|, elapsed_time, Z$ ", where X

is the number of edges representing route from customer to customer under a bus, T is the time window constraint of each customer, and Z is the objective value of the solution. Then is aggregated under pandas dataframe to be displayed in the report with matplotlib.

The experiment is conditioned on two variants, no limit on number of vehicles being routed by adjusting the constraint 5.3 and 5.5, number of vehicles out and in from depot to be at most 1, and the second variant is where we cheat by setting the number of vehicle to be the number of the optimal solution to gain insights and reduce run time by cutting searching space.

Results

Firstly, the computation time for the first variant, where no limit on number of vehicle is set, breaks CPLEX at 25 customers node specifically at C102 and beyond, thus the second variant will be the focal point of the experiment.

25 Customer Solution Quality

CPLEX is able to solve almost all of the instances of 25 customers, with the exception of R110 and R111, where R110 is the instance with no solution, but R111 was terminated by hand. Below is the table of comparison between the known optimal solution and CPLEX solution.

Table 1: Percentage gap between CPLEX results and benchmark optimal solutions for Solomon 25 customer

Instance	Percentage Gap
C101	0.268489
C102	0.229968
C103	0.229968
C104	0.293984
C105	0.268489
C106	0.268489
C107	0.268489
C108	0.268489
C109	0.268489
C201	0.392438
C202	0.392438
C203	0.392438
C204	0.389521
C205	0.392438
C206	0.392438
C207	0.389761
C208	0.407234
R101	0.199306
R102	0.184211
R103	0.241584
R104	0.254552
R105	0.195785
R106	0.232104
R107	0.228264
R108	0.250367
R109	0.300295
R111	12.063059
R112	0.532028
R201	0.231965
R202	0.240534
R203	0.236379
R204	0.251148
R205	0.270233
R206	0.289378
R207	0.286116
R208	0.345369
R209	0.232497
R210	0.217092
R211	0.288048
RC101	0.229006
RC102	0.268411
RC103	0.336129
RC104	0.175800
RC105	0.261752
RC106	0.291012
RC107	0.217848
RC108	0.167891
RC201	0.289013
RC202	0.243755
RC203	0.242658
RC204	0.178364
RC205	0.274971
RC206	0.340873
RC207	0.217848
RC208	0.288669

The percentage gap have mean of 1.5906213476904933% and standard deviation of 0.4928704901050774%, without R111, the mean and std becomes 0.07213788788994965%, 0.2786077471971899%. The formulation is very effective in terms of accuracy.

25 Customer Solution runtime

Based on 25 customers data set generally Clustered dataset has the best performance, with shortest runtime followed by random clustered dataset, and random dataset having the longest runtime.

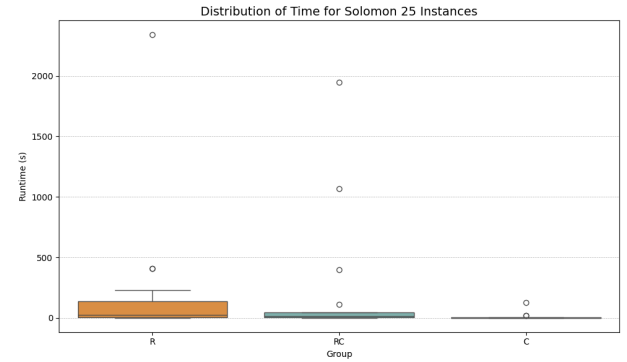


Figure 2: distribution of runtime for 25 customers in seconds

50 and 100 Customers

current formulation is not able to solve most 50 and 100 customers dataset, breaking in most cases except for C101 for 50 and 100 customers, and C102 for 50 customers. With so many little data points, it is hard to draw any conclusion on both solution quality and runtime.

Variable and Constraint Growth




Figure 3: Average number of variables for each dataset, at 25, 50 and 100 customers




Figure 4: Average number of constraint for each dataset, at 25, 50 and 100 customers

The dataset somewhat show explosive growth of the variables and constraints as it is shown to jump from 10k to 100k decision variables within 100 customers

For dataset of 50 customers, the runtime is sub one second for the first cluster set, then for the rest the program freezes and Cplex does not terminate. For dataset of 100 customers, CPLEX crashed and burned due to running out of memory, so experiments beyond 50 customers calls for better algorithm/ formulation to solve the problem.

Analysis and Discussion

As shown in the results section, the accuracy of the model is very high except for one instance, it is my opinion that this is due to numerical rounding instead of the model itself, as 0.3% is very small percentage gap to be considered another solution. Ideally the solution could be validated, but the optimal solution is not available for us. Aside from that, it is clear that Vehicle Routing Problem with Time Windows (VRPTW) is computationally challenging, especially as problem size increases. A key issue is the rapid growth of decision variables—particularly the binary arc variables x_{ijk} —which scale cubically with the number of customer nodes and number of vehicles. Specifically, the number of edges is $O(|V|^2)$, and with each vehicle k , the number of variables grows to $O(|V|^2 \cdot |K|)$. Although the number of vehicles $|K|$ is typically fixed (e.g., around 10 in the Solomon dataset), this growth, reflected in the figures above, still leads to an explosive number of variables. and since each variable is binary integer. variable, they are subjected to branch-and-bound search, which further increases computational expense.

The runtime result for 25 customers show runtime varies significantly by dataset structure with the clustered dataset achieving the shortest runtimes, random clustered datasets following, and purely random datasets exhibiting the longest runtimes.

This is to be expected, since clustered dataset has more limited search space, and thus less spaces to perform branch and bound on, while random dataset has the more space to branch and bound on, thus the longest runtime. It is to be noted however that the distribution of runtime is very random due to multiple factors, random input, CPLEX’s internal heuristics, so this distribution is not the representation of the actual runtime, but rather a representation of the runtime of the solver on the specific dataset. Generally the trend from experiment is that the runtime is either sub one second, or 2000 seconds, or does not terminate at all, the extreme cases make it hard to draw any conclusion on the runtime performance of the solver. The fact that distance can play a role in runtime in unknown ways make it harder to draw any conclusion of what causes CPLEX to break. Distance plays a role but its effect remains unclear.

Overall, the results suggest that exact methods such as branch-and-bound are not practical for solving large VRPTW instances. Even a state-of-the-art solver like CPLEX struggles under default configurations. This underlines the need for more scalable solution techniques, such as heuristics and metaheuristics—e.g., genetic algorithms, simulated annealing, or ant colony optimization—which are better suited to navigating the vast combinatorial search space of VRPTW in reasonable time at the cost of optimality.

Conclusion

This report presented a mathematical formulation and computational study of the Vehicle Routing Problem with Time Windows (VRPTW) using a mixed-integer programming (MIP) approach. The formulation captures both capacity and time window constraints, with the primary objective of minimizing total Euclidean travel distance. Assumptions included an unlimited fleet size, no penalty for waiting time, and valid preprocessing conditions.

Computational experiments conducted on Solomon benchmark datasets revealed that the model is capable of producing near-optimal solutions for instances with up to 25 customers. However, as the instance size increases, the model exhibits exponential growth in the number of binary decision variables—scaling as $O(|V|^2 \cdot |K|)$ —and consequently suffers from significant performance degradation in terms of both runtime and memory usage. In particular, CPLEX was unable to solve several instances beyond 25 customers under default configurations.

The analysis further indicates that problem structure significantly affects solver performance: clus-

tered customer distributions lead to substantially faster convergence than random or mixed distributions, due to a more constrained and localized feasible region.

In summary, while the MIP formulation provides a rigorous framework for modeling VRPTW, its scalability and time-feasibility is limited. These findings highlights the need for alternative solution techniques such as better decomposition methods and heuristics to tackle larger instances effectively.

References

- [1] Bernabé Dorronsoro. Vrp benchmarks: Solomon instances, 2007. Accessed: 2025-07-15.
- [2] Victor Pillac, Michel Gendreau, Christian Guéret, and Andrés L. Medaglia. The cvrp with time windows (cvrptw). https://www.researchgate.net/figure/The-CVRP-with-Time-Windows-CVRPTW_fig3_319754352, 2013. Accessed: 2025-07-15.
- [3] Paolo Toth and Daniele Vigo, editors. *Vehicle Routing: Problems, Methods, and Applications*. Society for Industrial and Applied Mathematics and Mathematical Optimization Society, Philadelphia, 2nd edition, 2014.